國立臺灣大學電機資訊學院電機工程學系
碩士論文
Department of Electrical Engineering
College of Electrical Engineering and Computer Science
National Taiwan University
Master Thesis

利用結構性支撐向量機的具音樂表現能力之半自動電腦演奏
系統
A Semi-automatic Computer Expressive Music Performance
System Using Structural Support Vector Machine

呂　行
Shing Hermes Lyu

指導教授：鄭士康博士
Advisor: Shyh-Kang Jeng, Ph.D.

中華民國 103 年 6 月
June, 2014

國立臺灣大學
電機工程學系

碩士論文

利用結構性支撐向量機的具音樂表現能力之半自動電腦演奏系統

呂 行 撰

103
6

國立臺灣大學碩士學位論文
口試委員會審定書

利用結構性支撐向量機的具音樂表現能力之
半自動電腦演奏系統
A Semi-automatic
Computer Expressive Music Performance System
Using Structural Support Vector Machine

　　本論文係呂行君（學號 R01921032）在國立臺灣大學電機工程學系完成之碩士學位論文，於民國 103 年 6 月 5 日承下列考試委員審查通過及口試及格，特此證明。

口試委員：

_____（簽名）
　　　　　　　　　　　　　　　　　　　（指導教授）

_____　　_____

_____　　_____

_____　　_____

_____　　_____

系主任　　_____（簽名）

誌謝

中文摘要

　　電腦合成的音樂一向被認為是僵硬、機械化而且沒有音樂表現能力。因此能夠產生具有表現能力的電腦自動演奏系統將會對音樂產業、個人化娛樂以及表驗藝術領域有重大的影響。在這篇論文中，我們藉由隱藏式馬可夫模型結構的結構性支撐向量機 (SVM-HMM) 來設計一個可以產生具有表現能力音樂的電腦自動演奏系統。我們邀請六位研究生錄製了克萊門蒂（Muzio Clementi）的小奏鳴曲集 Op.36。我們手動將這些錄音分割成樂句，並且利用程式從中抽取出音樂特徵。這些音樂特徵藉由 SVM-HMM 訓練成數學模型後，可以利用這個數學模型來演奏訓練過程中沒有見過的樂譜（需要手動標注樂句）。此系統目前只能支援單音旋律。問卷調查的結果顯示，本系統產生的音樂尚不能達到真人的演奏水準。但是根據量化的相似度分析，本系統產生的音樂確實比無表現性的 MIDI 音樂更接近真人演奏。

　　關鍵字：電腦自動演奏、結構性支撐向量機、支撐向量機

Abstract

Computer generated music is known to be robotic and inexpressive. A computer system that can generate expressive performance can potentially have significant impact on music production industry, personalized entertainment or even art. In this paper, we have designed and implemented a system that can generate expressive performance using structural support vector machine with hidden Markov model output (SVM-HMM). We recorded six sets of Muzio Clementi's Sonatina Op.36 performed by six graduate students. The recordings and scores are manually split into phrases and had their musical features automatically extracted. Using the SVM-HMM algorithm, a mathematical model of expressive performance knowledge is learned from these features. The trained model can generate expressive performances for previously unseen scores (with user-assigned phrasing). The system currently supports monophonic music only. Subjective test shows that the computer generated performance still can achieve the same level of expressiveness of human perfromers, but quantitative similarity measure shows that the computer generated performance is much similar to huamn performance than inexpressive MIDI.

Key words: Computer Expressive Performance, Performance Rendering, Structural SVMs, Support Vector Machines.

Table of Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Motivation

From the mechanical music performing automata from middle ages, to the latest Japanese virtual signer Hatune Miku, there had been many attempts to create automated systems that perform music. However, many of these systems can only generate predefined expression. State-of-the-art text-to-speech system can already generate fluid and natural speech, but computer performance still can't perform very expressively. Therefore, many researcher have devoted their effort to develop systems that can automatically or semi-automatically perform music expressively. There is even a biannual contest for such systems called Music Performance Rendering Contest (RenCon) [1]. The RenCon roadmap suggest that by 2050, they wish that a computer performer can win the Chopin International Piano Contest.

There are many potential applications for a computer expressive performance system, many commercial music typesetting software like Finale [2] and Sibelius [3] already have expressive playback features built-in. For entertainment industry, such system can provide personalized music listening experience. For music production industry, this technology can save a lot of cost on hiring musicians and license fees. Such system also opens up new opportunity in art, such as human-machine co-performance or interactive multimedia installation. In academia, researchers can use this technology to study the performance style of musicians, or restore historical recording archive.

1.2   Goal and Contribution

The ultimate goal of this paper is to be able to play any music in any expressive style specified. But due to technical and time constrains, we narrow down our goal to building a computer expressive performance system that performs monophonic music phrases by off-line supervised learning. The phrasing need to be annotated by human, so it's a semi-automatic system.

The major contribution of this paper is that we apply structural support vector machine on expressive performance problem. There exist no previous work that uses the discriminative learning power of structural support vector machine with hidden Markov model output (SVM-HMM) on computer expressive performance question. We also developed methods and tools to generate a expressive performance corpus.

1.3   Chapter Organization

In Chapter 2, we will give an overview of previous works and their varying goals, these works will be grouped by the way they learn performance knowledge, and we will discuss some additional specialities such as special instrument model or special user interaction pattern. In Chapter 3, we will first give a brief introduction to the mathematical background of SVM-HMM, and then give a top-down explanation to the proposed method. In Chapter 4, we will explain how the corpus used for training is designed and implemented. In Chapter 5, we will discuss several experiments that demostrates design trade-offs and the subjective test results. Finally, we have included an appendix that presents some software tools used in this research, which may be helpful for other researchers in the computer music field.

Chapter 2

Previous Works

2.1   Various Goals and Evaluation

The general goal of a computer expressive performance system is to generate expressive music, as opposed to the robotic and dull expression of rendered MIDI. But the definition of "expressive" is very vague and ambiguous, so each research will need to define a more precise and measurable goal. The following are the most popular goals a computer expressive performance system wants to achieve:

1. Perform music notations in a non-robotic way (no specific style).

2. Reproduce a human performance or a certain musician's style.

3. Accompany a human performance.

4. Validate a musicological theory of expressive performance.

5. Directly render computer-composed music works.

Some systems try to perform music notations in a non-robotic way in a general sense, without a certain style in mind. These systems has been employed in music typesetting softwares, like Finale [2] and Sibelius [3], to play the notation expressively. Most systems will implicitly include this goal.

Systems that are designed to reproduce certain human performance or style are usually designed and trained using a particular performer's recordings. One commercial example

is the Zenph re-performance CD [4]. This CD contains music performed by an expressive performance model of Rachimaninov's style, but Rachmaninov had never recorded these pieces in his lifetime.

Accompaniment systems try to render expressive music that act as an accompaniment for a human performance. The challenge is that the system must be able to track the progress of a human performance and adaptively render the accompaniment in real-time. One commercial example is Cadenza [5], using the technology created by Christopher Raphel. It can track the soloist's performance and play the accompaniment orchestral part accordingly.

Another goal is to validate musicological theories. Musicologist may propose theories on expressive music performance, by building a generative model, they can validate their theories. These systems may focus more on the specific phenomenon that the theory tries to explain instead of generating music that is pleasant to human.

Finally, some systems combines computer composition with expressive performance. These systems have a big advantage because the intention of the composer can be shared with the performer. Other systems that performs past compositions can only guess the composer's intention by analyzing the score notation. These systems usually has their own data structure to represent music, which can contain more information than traditional music notation, but the performance system is not backward compatible with past compositions.

Because of the high diversity in the goals they want to achieve, it is very hard to make fair comparison between systems. But we can still evaluate the capability of these systems by the following three key indicators proposed by [6]:

1. Expressive expression capability

2. Polyphonic capability

3. Performance creativity

Expressive expression capability can range from very high level structural expression (e.g. tempo contrast between sections) to note level expression (e.g. onset, loudness,

duration) or even sub-note expression (e.g. loudness envelop, timbre). Most systems can generate note-level expression, but higher or lower level expressions are much rare.

Polyphonic capability indicates if the system can perform polyphonic input. Polyphonic systems are more challenging than monophonic ones because they requires synchronization between voices.

Performance creativity measures the ability of the system to create novel expression. The desired level of creativity varies from goal to goal. A system aiming to recreate human performance may want to produce deterministic expressions based on the learning material, while a system that is combined with a composition system may want to create highly novel performance.

Each system will design different experiment and metrics to verify their goals. Thus, the self-reported results are can hardly be compared. The only public contest that evaluates expressive performance systems is called RenCon (Performance Rendering Contest) [1]. Scores (MIDI) will be given to participants one hour before the competition starts. The participants must generate the expressive version of the MIDIs in the given time, the MIDIs will be played live on a Yamaha Disklavier piano. The audience and a jury cosists of professional musicians will give ratings for each performance. The performances are played in random order, so the audience and jury won't know which participant is behind each performance.

The RenCon is divided into fully automatic and semi-automatic categories. But the degree of human intervention in the semi-automatic category varies widely between systems. So it's not very fair to compare them.

2.2 Researches Classified by Methods Used

Despite the difference between goals of different expressive performance systems, all expressive performance systems must have some strategy to learn and apply performance knowledge. There are generally two approach: rule-based or machine-learning-based.

Using rules to generate expressive music is probably the earliest approach. Director Musices [7] is one of the early example. Pop-E [8] is also a rule-based system which

can generate polyphonic music, using its voice synchronization algorithm. Computational Music Emotion Rule System [9] tried to develop rules that express human emotions. Other systems like Hierarchical Parabola System [7, 10--12], Composer Pulse System [13, 14], Bach Fugue System [15], Trumpet Synthesis System [16, 17] and Rubato [18, 19] are also some examples. Most of the rule-based systems focus on expressive attributes like note onset, note duration and loudness, but Hermode Tuning System [20] put special emphasis on intonation. Rule-based systems are generally more computationally efficient because the mathematical model is much simple than those learned by machine learning algorithms. And rules are generally more understandable to human than complex model parameters. But some of the nuance, such as some subconscious deviation, may be hard to describe by rules, so there is a emperical limit on how complex the rule-based system can be. Lack of creativity is also a problem for rule-based approach.

Another approach is to acquire performance knowledge by machine learning. Many machine learning methods have already been applied to this problem. For example, Music Interpretation System [21--23] and CaRo [24--26] both use linear regression to learn performance knowledge. But it is very unlikely that the expressive performance problem is a linear system, so Music Interpretation System try to introduce non-linearity by using logic AND operations on linear regression results. But generally speaking, linear regression is too simple to capture the core of expressive performance.

More complicated machine-learning algorithms have also been applied: ANN Piano [27] and Emotional flute [28] uses artificial neural network. ESP Piano [29] and Music Plus One [30--32] uses statistical graphical models such as hidden Markov model (HMM) and Bayesian belief network, but they did no use structural support vector machine to train the HMM. KCCA Piano System [33] uses kernel regression. Drumming System [34] tried different mapping models that generates drum patterns.

Evolutionary computation such as genetic programming is used in Genetic Programming Jazz Sax [35], Sequential Covering Algorithm Genetic Algorith [36], Generative Performance Genetic Algorithm [37] and Multi-Agent System with Imitation [38, 39]. Evolutionary computation takes long training time, and the results are less predictable.

But being unpredictable also means that these systems can create interesting performances in an unconventional way.

Another possible approach is to use case-based reasoning. SaxE [40--42] use fuzzy rules based on emotions to generate Jazz saxophone performance. Kagurame [43, 44] focus on style (Baroque, Romantic, Classic etc.) instead of emotion. Ha-Hi-Hun [45] has a more ambitions goal in mind: to accept natural language instructions like "Perform piece X in the style of Y." Another series of researches done by Widmer at el., called PLCG [46--48], uses data mining technique to find rules for expressive performance. It's successor -- Phrase-decomposition/PLCG [49] added hierarchical phrase structures support to the original PLCG system. And the latest research in the series called DISTALL [50, 51] added hierarchical rules to the original one.

Most of the performance systems discussed above takes musical notation (MusicXML, MIDI, etc.) or inexpressive audio as input. They have to figures out the expressive intention of the composer by analyzing the score. But another type of computer expressive performance has a big advantage over the previous described ones, by combining computer composition and expressive performance, the performance module can receive the composition intention directly from the composition module. Ossia [52] and pMIMACS [53] are two examples of this category. This approach provides great possibility for creativity, but they can only play their own composition, which limits it range of application.

2.3   Additional Specialties

Most expressive performance systems implicitly or explicitly generate piano performance, because it's relatively easy to collect training samples for piano, and piano sound is relatively easy to synthesize. Yet, some systems generate music in other instruments, such as saxophone [40--42], trumpet [16,17], flute [28] and drums [54]. These systems requires extra effort in creating instrument models in training, generation and synthesizing. Y.-H Kuo et al. [55] also propsed a way to re-synthsize individual notes into a performance with smooth timbre variation, but the work focus more on sub-note level timbre systhsis.

If not specified, most systems handles traditional western tonal music. However, most

saxophone-based work [40--42] generates Jazz music, because saxophone is an iconic instrument in Jazz performance. And the Drumming System [54] generates Brazilian drumming music.

Performing polyphonic music is much more challenging than monophonic music, because it requires synchronization between voices. Pop-E [8] use a synchronization mechanism to achieve polyphonic performance. Bach Fugue System [15] is created using the polyphonic rules in music theory about fugue, so it's inherently able to play polyphonic fugue. KCCA Piano System [33] can generate homophonic music -- an upper melody with an accompaniment -- which is common in piano music. Music Plus One [30--32] is a little bit different because it's a accompaniment system, it adapts non-expressive orchestral accompaniment track to user's performance.

Chapter 3

Proposed Method

3.1 Overview

The high-level architecture of the purposed system is shown in Fig. 3.1. The system has two phases, the upper half of the figure is the learning phase, the lower half is the performing phase. In the learning phase, score and expressive human recording pairs, split into phrases by human, are used as training examples for structural support vector machine with hidden Markov model output (SVM-HMM) algorithm to learn performance knowledge model. In the performing phase, a score will be given to the system for expressive performance. The SVM-HMM generation module will use the performance knowledge learned in the previous phase to produce expressive performance. The SVM-HMM output then go through a MIDI generator and MIDI synthesizer to produce audible performance.

All the scores and recordings are monophonic and contains only one musical phrase. The phrasing is done by human, thus the system is semi-automatic. The learning algorithm, namely SVM-HMM, can only perform off-line learning, so the learning phase can only work in a non-realtime scenario. The generating phase can work much faster, expressive music can be generated almost instantaneously.

There are many ways the user can control the performance style of the final output: first, the user can choose the training corpus. Theoratically, a model of a particular style can be learned from a set of samples with that particular style. Second, the user can control the structural expression by assigning the phrasing.

9

Figure 3.1: High-level system architecture

In the following sections, we will give an overview of the theroratical background behind SVM-HMM, and then walk through the detail steps in the learning and performing phases, and some implementation detail. The features used will be presented in the end of this chapter.

3.2   A Brief Introduction to SVM-HMM

In this thesis, we use structural support vector machine to learn performance knowledge from expressive performance samples. Unlike traditional SVM algorithm, which can only produce univariate prediction, structural SVM can produce structural predictions like tree, graph or sequence. Structural SVM with hidden Markov model output (SVM-HMM) has been successfully applied to part-of-speech tagging problem [56]. The part-of-speech tagging problem has some similarity with expressive performance problem. In part-of-speech tagging, one tries to identify the role by which the word plays in the sentence, while in expressive performance, one tries to determine how a note should be played, usually based on it's role in the musical phrase. For example, an authentic cadence at the end of a phrase is usually played louder and stronger than a embellishment note in the middle of a phrase. Thus, we believe SVM-HMM will be a good candidate for expressive performance. The following introduction and formulas relies heavily on [56--58].

10

Traditional SVM prediction problem can be described as finding a function

$$h : \mathcal{X} \rightarrow \mathcal{Y}$$

with lowest prediction error. $\mathcal{X}$ is the input features space, and $\mathcal{Y}$ is the prediction space. In traditional SVM, elements in $\mathcal{Y}$ are labels (classification) or real values (regression). But structural SVM extends the framework to generate structural output, such as tree, graph or sequence. To extend SVM to support structured output, the problem is modified as finding a discriminant function

$$F : \mathcal{X} \times \mathcal{Y} \rightarrow \mathcal{R}$$

, in which the input/output pairs are mapped to a real number score. To predict an output y for an input x, one try to maximize F over all $y \in \mathcal{Y}$.

$$f(x) = \arg\max_{y \in \mathcal{Y}} F(w, x, y)$$

Let F be a linear function of the form:

$$F = \mathbf{w}^T \Psi(x, y)$$

, where $\mathbf{w}$ is the parameter vector, and $\Psi(x, y)$ is the kernel function relating input x to output y. Ψ can be defined to accommodate various kind of structures.

For each structure we want to predict, a loss function that measures the accuracy of of a prediction is required. A loss function $\Delta : \mathcal{Y} \times \mathcal{Y} \rightarrow R$ need to satisfy the following property:

$$\Delta(y, y') \geq 0 \; for \; y \neq y'$$

$$\Delta(y, y) = 0$$

The loss function is assumed to be bounded. Let's assume the input-output pair (x, y) is drawn from a join distribution P(x,y), the prediction problem is to minimize the total

loss:

$$R_p^\Delta = \int_{\mathcal{X} \times \mathcal{Y}} \Delta(y, f(x)) dP(x, y)$$

Since we can't directly find the distribution P, we need to replace this total loss with a empirical loss, which can be calculated from the observed training set of (x_i, y_i) pairs.

$$R_s^\Delta(f) = \frac{1}{n} \sum_{i=1}^{n} \Delta(y_i, f(x_i))$$

Now we are ready to extend SVM to structural output, starting with a linear separable case, and we will then extend it to soft-margin formulation.

A linear separable case can be expressed by a set of linear constrains

$$\forall i \in \{1, \cdots, n\}, \forall \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i, y_i) - \Psi(x_i, \hat{y}_i)] \geq 0$$

The constrains imply that the groundtruth y_i for x_i has the minimum F value than any other $\hat{y}_i \neq y_i$.

The key concept of SVM is the large margin principle. We not only want to find a solution that statisfies the constrains, but also we want to maximize the margin between the groundtruth and the second best $\hat{y}_i$:

$$\max_{\gamma, \mathbf{w}: \|\mathbf{w}\|=1} \gamma$$
$$s.t \ \forall i \in \{1, \cdots, n\}, \forall \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i, y_i) - \Psi(x_i, \hat{y}_i)] \geq \gamma$$

, which is equivalent to the convex quadratic programming problem:

$$\min_{\mathbf{w}, \xi_i \geq 0} \frac{1}{2} \|\mathbf{w}\|^2$$
$$s.t. \ \forall i \in \{1, \cdots, n\}, \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i, y_i) - \Psi(x_i, \hat{y}_i)] \geq 1$$

To extend the linear-separable case to non-separable case, slack variables ξ_i can be

introduced to penalize prediction errors, results in a soft-margin formalization:

$$\min_{\mathbf{w},\xi_i \geq 0} \frac{1}{2}\|\mathbf{w}\|^2 + \frac{C}{n}\sum_{i=1}^{n}\xi_i$$

$$s.t. \ \forall i \in \{1,\cdots,n\}, \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i,y_i) - \Psi(x_i,\hat{y}_i)] \geq 1 - \xi_i$$

C is the weighting parameter controlling the trade-off between low training error and large margin. The optimal C varies between different problems, so experiment should be conducted to find the optimal C for our problem.

Intuitively, a constrain violation with larger loss should be penalize more than the one with smaller loss. So I. Tsochantaridis et al. [57] proposed two possible way to take the loss function into account. The first way is to re-scale the slack variable by the inverse of the loss, so a high loss leads to smaller re-scaled slack variable:

$$\min_{\mathbf{w},\xi_i \geq 0} \frac{1}{2}\|\mathbf{w}\|^2 + \frac{C}{n}\sum_{i=1}^{n}\xi_i$$

$$s.t. \ \forall i \in \{1,\cdots,n\}, \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i,y_i) - \Psi(x_i,\hat{y}_i)] \geq 1 - \frac{\xi_i}{\Delta(y_i,\hat{y}_i)}$$

The second way is to re-scale the margin, which yields

$$\min_{\mathbf{w},\xi_i \geq 0} \frac{1}{2}\|\mathbf{w}\|^2 + \frac{C}{n}\sum_{i=1}^{n}\xi_i$$

$$s.t. \ \forall i \in \{1,\cdots,n\}, \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i,y_i) - \Psi(x_i,\hat{y}_i)] \geq \Delta(y_i,\hat{y}_i) - \xi_i$$

In the implementaion we use, we use margin re-scaling.

But the above quadratic programming problem has a very large number ($O(n|\mathcal{Y}|)$) of constrains, which will take considerable time to solve. I. Tsochantaridis et al. [57] proposed a greedy algorithm to speed up the process by selecting only part of the constrains that contributes the most to finding the solution. Initially, the solver starts with an empty working set containing no constrains. Than the solver iteratively scans the training set to find the most violated constrains under the current solution. If a constrain is violated more times than a desired threshold, the constrain is added to the working set of constrains.

Then the solver re-calculate the solution under the new working set. The algorithm will terminate once no more constrain can be added under the desired precision.

In a later work by Joachims et al. [56], they created a new formulation and algorithm to further speed up the algorithm. Instead of using one slack variables for each training sample, which results in a total of n slack variables, they use a single slack variable for all n training samples. The following formula is the 1-slack version of slack-rescaling structural SVM:

$$\min_{\mathbf{w},\xi_i \geq 0} \frac{1}{2}\|\mathbf{w}\|^2 + C\xi$$

$$s.t. \ \forall i \in \{1,\cdots,n\}, \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i,y_i) - \Psi(x_i,\hat{y}_i)] \geq \frac{1}{n}\sum_{i=1}^{n} 1 - \frac{\xi}{\Delta(y_i,\hat{y}_i)}$$

And margin-rescaling structural SVM:

$$\min_{\mathbf{w},\xi_i \geq 0} \frac{1}{2}\|\mathbf{w}\|^2 + C\xi$$

$$s.t. \ \forall i \in \{1,\cdots,n\}, \hat{y}_i \in \mathcal{Y} : \mathbf{w}^T[\Psi(x_i,y_i) - \Psi(x_i,\hat{y}_i)] \geq \frac{1}{n}\sum_{i=1}^{n} \Delta(y_i,\hat{y}_i) - \xi$$

Detailed proof on how the new formulation is equally general as the old one is given in the paper [56].

With the framework described above, the only problem left is how to define the general loss function and Ψ for our problem? Drawing the inter-state dependencies and time dependencies concept from hidden Markov model, Y. Altun et al. [58] proposed two types of features for a equal-length observation/label sequence pair (x,y). The first is the interaction of a observed feature x^s with a label y^t, the other is the interaction between neighboring labels y^s and y^t.

To illustrate the method, we use a example from music: for some observed features $\Psi_r(x^s)$ of a note x located in s-th position of the phrase, and assume $[[y^t = \tau]]$ denotes the t-th note is played at a velocity of τ, the interaction of the observed feature and the label

can be written as:

$$\psi_{r\sigma}^{st}(\mathbf{x}, \mathbf{y}) = \left[\!\left[y^t = \tau \right]\!\right] \Psi_r(x^s), \ 1 \leq \gamma \leq d, \ \tau \in \Sigma$$

And the interaction between labels can be written as:

$$\hat{\psi}_{r\sigma}^{st}(\mathbf{x}, \mathbf{y}) = \left[\!\left[y^s = \sigma \wedge y^t = \tau \right]\!\right], \ \sigma, \tau \in \Sigma$$

By selecting a order of dependency for the HMM model, we can further restrict s's and t's. For example, for a first-order HMM, $s = t$ for the first feature, and $s = t - 1$ for the second feature. The two features on the same time t is then stacked into a vector $\Psi(x, y; t)$. The feature map for the whole sequence is simply the sum of all the feature vectors

$$\Psi(\mathbf{x}, \mathbf{y}) = \sum_{t=1}^{T} \Psi(\mathbf{x}, \mathbf{y}; t)$$

The distance, i.e. the general loss function, between two feature maps depends on the number of common label segments and the inner product between the input features sequence with common labels.

$$\Delta(\Psi(\mathbf{x}, \mathbf{y}), \Psi(\hat{\mathbf{x}}, \hat{\mathbf{y}})) = \sum_{s,t} \left[\!\left[y^{s-1} = \hat{y}^{t-1} \wedge y^s = \hat{y}^t \right]\!\right] + \sum_{s,t} \left[\!\left[y^s = \hat{y}^t \right]\!\right] k(x^s, \hat{x}^t)$$

Finally, during the prediction process, a Viterbi-like decoding algorithm is used to effeciently find a y that maximize F.

3.3   Learning Performance Knowledge

In this section, we will introduce the componants that consist the learning phase. The main goal in the learning phase is to extract performance knowledge from training samples. Fig. 3.2 shows the internal structure of the learning phase.

Figure 3.2: Learning phase flow chart

Training samples are matched score and expressive performance pairs (their format and preparation process is discussed in Chapter 4). The raw data from the samples are too complex to process, so we need to extract important features from them. Two types of features will be extracted from the samples: first, the musicological cues from the scores are called score features; second, the measurable expression from the expressive performances are called the performance features. We want the system to learn how score features are "translated" into performeance features. This process can be analogize to a human performer reading the explicit and implicit cues from the score, and perform the music with certain expressive expression. The definition of the features used will be presented in Section 3.5.

3.3.1 Training Sample Loader

The training samples are loaded by the sample loader module. Since a training sample consists of a score (musicXML format) and an expressive recording (MIDI format), the sample loader finds the two files, and load them into an intermediate representation (`music21.Stream` object provided by the `music21` library [59] from MIT). The music21 library will convert the musicXML and MIDI format into a Python Object hierarchy that is easy to access and manipulate by Python code.

One caveat here is the music21 library will quantize the time in MIDI, which will destroy the subtle onset and duration expressions. And the music21 library don't handle the "ticks per quarter note" information in the MIDI header [60], which is essential for the MIDI parser to interpret the correct time scale. So we must explicitly disable quantization and specify the "ticks per quarter note" value during MIDI loading.

3.3.2 Features Extraction

In order to keep the system architecture simple, feature extractors are designed to be independent to other feature extractors, so features can included or removed without affecting the rest of the system. Furthermore, this enables parallel feature extraction. But sometimes a feature inevitably depends on other features, for example, the "relative duration with the previous note" is calculated based on the "duration" feature. Since we want to avoid complex dependency management, the "relative duration with the previous note" feature extractor has to invoke the "duration" extractor, instead of waiting for the "duration" extractor to finish first. Therefore, the "duration" feature extracted will be computed twice. To avoid redundant computation of the feature extractors, we implemented a caching mechanism. Once the "duration" feature had been computed, no matter it is calculated during "duration" extraction or or during "relative dutaion with the previous note" extraction process, it's value will be cached during this execution session. So no matter how many feature extractors uses the "duration" feature, they can get the value directly from cache. This method can speed up the execution without needing to handling dependencies.

The extracted features are aggregated and stored into a JavaScript Object Notation (JSON) file for the SVM-HMM module to load. By saving the features in a human-readable intermediate file, we can debug potential problems easily.

3.3.3 SVM-HMM Learning

After all features are extracted, the next step is to learn performance knowledge from the features. In the early stage of this research, we have successfully applied linear regressio [61]. However, assuming this problem to be linear is clearly an oversimplification, so we switch to structural support vector machine with hidden Markov model output (SVM-HMM) [56--58] as our supervised learning algorithm.

The SVM-HMM learning module loads the feature file from the previous stage, and aggregate the features to fit the required input format of the SVM-HMM learner program. However, most features from the previous stage are real values, but SVM-HMM only takes

discrete performance features[1], so quantization is required. There are many possible way to quantize the features, each will result in different output, here we will present a quantizer design for demonstration purpose: for each performance feature, the mean and standard deviation from all training samples are calculated first. The range between mean minus or plus four standard deviations is divided into 128 uniform intervals. Values over than mean plus four standard deviations are quantized into the 128th bin, and values below mean minus four standard deviations are quantized into the 1st bin. The number of intervals decides how fine-grain the quantization is, if the number is too low, subtle expressions will be lost due to high quantization error. However, if the number is too large, there will be too few samples for each interval, which is bad from a statistical learning perspective. Also the training process will take a lot of CPU and memory resources without significant gain in prediction accuracy. The range of four standard deviation is chosen by trail and error, a narrower range will make most of the extreme values be quantized into the largest of smallest bin, so the performance will have a lot of saturated values. But a very large range will make the interval between each quantization bin too large, rising the quantization error.

The theoretical background of SVM-HMM is already mentioned in Section 3.2. We leverage Thorsten Joachims's implementation called SVM^{hmm} [62]. SVM^{hmm} is an implementation of structural SVMs for sequence tagging [58] using the training algorithm described in [57] and [56]. The SVM^{hmm} package contains a SVM-HMM training program called `svm_hmm_learn` and a prediction program called `svm_hmm_classify`. For architectural simplicity, we train one model for each performance feature, each model uses all the score features to predict a single performance feature. The `svm_hmm_learn` read the features from a file in the following format: Each line represents features for a note in time order, format as

```
PERF qid:EXNUM FEAT1:FEAT1_VAL FEAT2:FEAT2_VAL ... #comment
```

`PERF` is a quantized performance feature. The `EXNUM` after `qid:` identifies the phrases, all notes in a phrase will have the same `qid:EXNUM` identifier. Following the identi-

---

[1]SVM-HMM is initially designed for tasks like part-of-speech tagging, in which real value or binary features are used to predict discrete part-of-speech tags.

18

fier are quantized score features, denote as `feature name : feature value`, separated by spaces. And text following a # symbol is comment.

There are some key parameters needed to be adjusted for the training program. First the C parameter in SVM, which controls the trade-off between lowering training error and maximizing margin. Larger C will result in lower training error, but the margin may be smaller. Second, the ε parameter controls the required precision for termination. The smaller the ε, the higher precision, but it may require more time and computing resource. Finally, for the HMM part of the model, the order of dependencies of transition states and emission states needs to be specified. In our case, both are set to defaults: transition dependency is set to one, which stands for first-order Markov property, and emission dependency is set to zero. Since we train one models for each performance feature, each model can have their own set of parameters. The parameter selection experiments will be presented in Chapter 5.

Finally, the training program will output three model files (because we use three performance features) which contains SVM-HMM model parameters, such as the support vectors and other metadata. Since it takes considerable time (roughly a dozen minutes to a few hours) to train a model, depending on the amount of training samples and the power of the computer, the system can only support off-line learning. But the learning process only need to be run once. The performance knowledge model can be reused over and over again in the performing phase.

3.4   Performing Expressively

The performing phase uses the performance knowledge model learned in the previous phase to generate expressive performances. The input is a score file to be performed, which should not be used as training sample to prevent overfitting. Score features will be extracted from it using the same routine as in the learning phase. The SVM-HMM generation module will use the learned model and the score features to predict the performance features. These features will than be de-quantized back to real values using the method described previously. An MIDI generation module will apply those performance features

Figure 3.3: Performing phase flow chart

onto the score to produce a expressive MIDI file. The MIDI file itself is already a expressive performance, in order to listen to the sound, an software synthesizer can be used to render the MIDI file into WAV or MP3 format.

3.4.1 SVM-HMM Generation

The feature extraction and aggregation process in the performing phase is similar to the learning phase, but the `PERF` fields in the SVM-HMM input file are left blank for the algorithm to predict. The `svm_hmm_classify` program will take these inputs with the learned model file and predict the quantized labels of the performance features. These performance features are de-quantized back to the middle point of each bin.

3.4.2 MIDI Generation and Synthesis

The predicted performance features are then applied onto the input score, i.e. the onset timings will be shifted, the duration extended or shortened, and the loudness shifted according to the predicted performance features. The resulting expressive performance will be transfromed into MIDI files using `music21` library [59].

In order to actually hear the expressive performance, the MIDI file can be rendered by a software MIDI synthesizer. For example, `timidity++` software synthesizer for Linux can render the MIDI into a WAV (Waveform Audio Format) file, which can be compressed into MP3 (MPEG-2 Audio Layer III) by `lame` audio encoder. Alternatively, one can use

hardware synthesizers, for example, RenCon [1] contest uses Yamaha Disklavier digital piano to render contestants' submission.

Because sub-note level expression is not the primary goal of this research, we choose standard MIDI grand piano sound to render the music. The system can be extended to used more advanced physical model or instrument-specific audio synthesizer. Some sub-note level features, such as special techniques for violins, can be added to the features list and be learned by the SVM-HMM model.

3.5 Features

As mentioned in Section 3.3, there are two types of features, score features and performance features. We will present the features used in the system, and discuss the difficulties encountered.

3.5.1 Score Features

Score features are musicological cues presented in the score. The purpose of score features are to simulate the high level information a performer may perceive when he/she reads the score. The basic time unit for these features are notes. Each note will have all features presented below. Score features includes:

Relative position in the phrase: The relative position of a note in the phrase, its value ranges from 0% to 100%. This feature is intended to capture the special expression in the start or the end of a phrase, or time-variant expression like arch-type loudness variation.

Pitch: The pitch of a note denoted by MIDI pitch number (resolution is down to semitone).

Interval from the previous note: The interval between the current note and its previous note (in semitone). This represents the direction of the melodic line.See Fig. 3.4 for

Figure 3.4: Interval from/to neighbor notes

example.

$$\Delta P^- = P_i - P_{i-1}$$

Interval to the next note: The interval between the current note and its previous note (in semitone). See Fig. 3.4 for example.

$$\Delta P^+ = P_{i+1} - P_i$$

Note duration: The duration of a note (quarter notes).

Grace notes have no duration in musicXML specification [63]. The reason for this is that grace notes are considered very short ornaments that does not occupy real beat position. But zero duration is hard to handle in math formulation. So we assigned a duration of a sixty-fourth note, because it's far shorter than all the notes in our corpus.

Relative Duration with the previous note: The duration of a note divided by the duration of its previous note. See Fig. 3.5 for example. For a phrase of n notes with duration $D_1, D_2, \ldots, D_n$,

$$RD^- = \frac{D_i}{D_{i-1}}$$

This feature is intended to locate local changes in tempo, such as a series of rapid consecutive notes followed by a long note, which will cause a discontinuity in this feature.

Relative duration with the next note: The duration of a note divided by duration of its

Figure 3.5: Relative duration with the previous/next note

Figure 3.6: Metric position

next note. See Fig. 3.5 for example.

$$RD^+ = \frac{D_i}{D_{i+1}}$$

Metric position: The position (beat) of a note in a measure. For example, under a time signature of $\frac{4}{4}$, if a measure consists of five notes, they will have metric position of 1, 2, 2.5, 3 and 4, respectively.

Metric position usually implies beat strength. In most tonal music, there exist a hierarchy of beat strength. For example, for a time signature of $\frac{4}{4}$, the first note is usually the strongest, the third note is the second strongest, and the second and fourth notes are the least strong.

3.5.2 Performance Features

Performance features are the expressive expressions we would like to learn from a performance. Performance features are extracted by calculating how the expression deviates from the nominal notation in the score. Performance features includes:

Onset time deviation: A human performer usually adds conscious or unconsious rubato ot their performance. The onset time deviation is the difference of onset timing

23

Figure 3.7: Systematic bias in onset deviation

between the performance and the score. Namely,

$$\Delta O = O_i^{perf} - O_i^{score}$$

Where O_i^{perf} is the onset time of note i in the performance, O_i^{score} is the onset time of note i in the score.

However, the above formula assumes the performance is played at the exact same tempo assigned by the score. However, performers can't always keep up with the speed of the score because of limited piano skill, or they may speed up or slow down certain sections as expression. Therefore, the performance should be linearly scaled to avoid systematic bias, We will present a solution to this issue in Section 3.5.3.

Loudness: The loudness of a note. Measured by MIDI velocity level 0 to 127.

Relative duration: The performed duration of a note divided by the nominal duration on the score.

$$RD = \frac{D_i^{perf}}{D_i^{score}}$$

3.5.3   Normalizing Onset Deviation

In the previous section, we mentioned that the onset deviation feature will have problems when the performer did not play at the exact tempo indicated by the score. As illustrated in Fig. 3.7, if the performance is played slower than expected, the deviation will grow larger and larger over time same, and vice versa is it's played faster. The systematic bias caused by the difference in total duration will mix up with the local deviation, For a

24

long phrase, the onset deviation of the last notes can be as larger as a dozen quarter notes. These kind of extremely large values will be learned by the model and cause erroneous predictions. A note may be delayed for a few quarter notes, causing it the notes to be played in the wrong order.

In other words, the onset deviation actually contains two type of deviation: a global/ systematic deviation cause by the difference between performed and nominal tempo, and a local deviation cause by note-level expression. Since the intention of the onset deviation feature is to capture the note-level expression, the performance must be linearly scaled to cancel out the global deviation.

Initially, we tried two possible way of normalization:

1. Align the onset of the first notes, and align the onset of the last notes.

2. Align the onset of the first notes, and align the end (MIDI note-off event) of the last notes.

However, neither of the method can robustly eliminate extreme values. Therefore, we proposed an automated approach to find the best scaling ratio such that the normalized onset deviations in the performances fits best with those in the score. The measure of fitting is defined as the Euclidean distance between the normalized performance onset sequences and the score onset sequences, represented as vectors. Brent's Method [64] is used to find this optimal ratio. To speed up the optimization and prevent unreasonable local minima value, a search range of $[initial\ guess \times 0.5, initial\ guess \times 2]$ is imposed on the optimizer. The $initial\ guess$ is used as a rough estimate of the ratio, calculated by aligning the first and last onsets. Than we assume the actual ratio will not be smaller than half of $initial\ guess$ and not larger than twice of $initial\ guess$. The two numbers 0.5 and 2 are chosen by trail and error, and most of the empirical data supports this decision. We will demonstrate the effectiveness of this solution in Section 5.1.

Chapter 4

Corpus Preparation

An expressive performance corpus is a set of performance samples. Since this research is based on a supervised learning algorithm, a high-quality corpus is essential to our success. Each sample consists of a score and its corresponding human recording. Some metadata such as phrasing, structure analysis, or harmonic analysis. may also be included. In this chapter, we will review some the existing corpora, specifications and formats of our corpus, and how we actually construct it.

4.1   Existing Corpora

Unlike other research fields like speech processing or natural language processing, there exist virtually no public accessible corpus for computer expressive performance research. CrestMusePEDB [65] (PEDB stands for "Performance Expression Database") is a corpus created by Japan Science and Technology Agency's CREST program. However, until the time of this writing, we can't establish any contact with the database administrators to gain access to it. They claims to have a GUI tool for annotate the expressive performance parameters from audio recordings. Their repertoire covers many piano works from well-known classical composers like Bach, Mozart, and Chopin, and are recorded by world famous pianists. From their website [65] they claim to contain the following data: PEDB-SCR - score text information, PEDB-DEV - performance deviation data and PEDB-IDX - audio performance credit. But the quality of the data is unknown.

Another example is the Magaloff Project [66], which is created by some universities in Austria. They invited Russian pianist Nikita Magaloff to record all solo works for piano by Frederic Chopin on a Bösendorfer SE computer-controlled grand piano. This corpus became the material for many subsequent researches [67--73]. Flossmann et al., one of the leading researchers of the project, also won the 2008 RenCon contest with a system based on this corpus called YQX [74]. However, the corpus is not opened up to the public.

Since both corpora are not available, we need to implement our own . We will start by defining the specification.

4.2   Corpus Specification

The corpus we need must fulfill the following criterias:

1. All the samples are monophonic, containing only a single melody without chords.

2. No human error, such as insertion, deletion, or wrong pitch exist in the recording; the score and recording are matched note-to-note.

3. Phrasing is annotated by human.

4. The scores, recordings and phrasing data are in machine-readable format.

Some potentially useful information are not included because they are less relevant to our goal. Examples are:

1. Advanced structural analysis, such as GTTM (Generative Theory of Tonal Music) [75]

2. Harmonic analysis

3. Piano paddle usage

4. Piano fingering

5. Other instrument specific techniques, such as violin pizzicato, tapping, or bow techniques.

Table 4.1: Clementi's Sonatinas Op.36

Title	Movement	Time Signature
No.1 Sonatina in C major	I. Allegro	4/4
	II. Andante	3/4
	III. Vivace	3/8
No.2 Sonatina in G major	I. Allegretto	2/4
	II. Allegretto	3/4
	III. Allegro	3/8
No.3 Sonatina in C major	I. Spiritoso	4/4
	II. Un poco adagio	2/2
	III. Allegro	2/4
No.4 Sonatina in F major	I. Con spirito	3/4
	II. Andante con espressione	2/4
	III. Rondó: Allegro vivace	2/4
No.5 Sonatina in G major	I. Presto	2/2
	II. Allegretto moderato	3/8
	III. Rondó: Allegro molto	2/4
No.6 Sonatina in D major	I. Allegro con spirito	4/4
	II. Allegretto	6/8

We choose Clementi's Sonatina Op.36 for our corpus, because it is a must-learn reper-toire for piano students, so it's easy to find performers with a wide range of skill level to record the corpus. These sonatinas are in classical style, so the learned model can easily be extended to other classical era works like Mozart and Haydn. There are six sonatinas included in Op.36, the first five have three movements each, and the last one has two movements. The titles and time signatures of all the pieces are listed in Table 4.1

MusicXML is used to represent Clementi's work in digital format. MusicXML is a digital score notation using XML (eXtensible Markup Language), it can express most traditional music notations and metadata. Most music notation software and software tool supports musicXML format. Although MIDI is also a possible candidate for representing score, it is designed to hold instrument control signal rather than notation, so some music symbols may not be available in MIDI. Furthermore, MIDI represents music as a series of note-on and note-off events, which requires additional effort to transform into traditional notation.

But for representing performance, MIDI is the most suitable format. Using a key-pressure-sensitive digital piano, pianist can record in a natural way. The recordings will have high precision in time, pitch and loudness (key pressure), and polyphonic tracks can

easily be recorded separately. Although WAV (Waveform Audio Format) audio recording has higher fidelity than MIDI, but they are harder to parse by computers. Without robust onset detection, pitch detection, and source separation technology, the information is extremely difficult to extract. Manually annotate each WAV recording takes unrealistic effort, and the accuracy across different annotators may not be consistent.

There is a way to keep both the score and the recording in one single MIDI file. Instead of recording the actual note-on and note-off timing, we keep the nominal note-on and note-off the same as in the score. Then, MIDI tempo-change events are inserted before each note to shift the performed timing of the recorded notes. Thus, the nominal time of each note represents the score, and the rendered time represents the performance. But since MIDI is so limited as a score format, and it requires complex calculations to recover the performance, this method is not used in the research.

Finally, we store the phrasing, which is the only metadata we used, in a plaintext file, each line in the phrasing file is the starting point of each phrase. The starting point is defined as the onset timing (in quarter notes) counted from the beginning of the piece[1]. The phrasing is decided by the us using the following principles:

1. Phrase may be separated by a salient pause.

2. Phrase may end with a cadence.

3. Phrase may be separated by dramatic change in tempo, key or loudness.

4. Repeated structures in tempo or pitch may be a repeated phrase.

Since phrasing controls the structural interpretation of a piece, we would like to leave this freedom for expression to the user. However, if there exist any good automatic phrasing algorithm, it can be easily integrated into the current system to make it full-automatic.

---

[1]For a phrase that start at a point which is a circulating decimal, for example $2\frac{1}{3} = 2.333\cdots$, the starting point can be alternatively defined as any finite decimal between the end of the last phrase and the start of the current phrase. For example, if the last phrase stops at beat 1, the second phrase start at $2\frac{1}{3} = 2.333\cdots$ beat, the start point of the second phrase can be written as 2.3 or 2.0, etc.

4.3 Implementation

4.3.1 Score Preparation

The digital scores are downloaded from KernScore website [76]. The scores are transformed into MusicXML from the original Hundrum file format (.krn) using the music21 toolkit [59]. Because this research focus on monophonic melody only, the accompaniments are remove and the chords are reduced to their highest-pitched note, which is usually the most salient melody. The reduced scores are doubled-checked against a printed version publish by Durand & Cie., Paris [77] to eliminate all errors.

4.3.2 MIDI Recording

We have implemented two methods for recording: First, using a Yamaha digital piano to record MIDI. Second, by tapping on a touch-sensitive device to express tempo, duration and loudness. Due to accuracy consideration, only the recordings from Yamaha digital piano are used in the expreiments.

We used a Yamaha P80 88-key graded hammer effect[2]digital piano for recording. Through a MIDI-to-USB converter, the keyboard was connected to Rosegarden Digital Audio Workstation (DAW) software on a Linux computer. The Rosegarden DAW also generated the metronome sound to help the performer maintain a steady speed. Metronome is mandatory because if the performer plays freely, the tempo information written in the MIDI file will be invalid, which makes subsequent parsing and linear scaling very difficult. So the performers were asked to follow the speed of the metronome, but they can adjust the metronome speed as they like, and apply any level of rubato as long as the overall tempo is steady.

The second method, which is not used in the experiments, is to utilize touch-enabled input device like smartphone touchscreen or laptop touchpad. We have implemented an prototype using a Synaptics Touchpad on a Lenovo ThinkPad X200i laptop. When the user taps the touchpad once, one note from the score will be played, the duration and loudness

---

[2]Graded Hammer Effect feature provides realistic key pressure response similar to a traditional acoustic piano.

will be controlled by the duration and pressure of the tapping action. So the user can "play" the touchpad like a musical instrument. This idea has already be used in musical games and toys. This method is more user-friendly to general public because it requires minimal instrument skill and utilize widely available hardware. But most touchpad estimates pressure by finger contact area, so the accuracy in pressure is not very satisfying. But it is indeed a low cost alternative to MIDI digital piano.

4.3.3   MIDI Cleaning and Phrase Splitting

After MIDIs are recorded, we use Python scripts to check if each recording is matched note-to-note with its corresponding score. If not, the mistakes are manually corrected. If there are a small segments that are totally messed up, they will be reconstruct using repeated or similar segments from the same piece. The matched score and MIDI pairs are then split into phrases according to the corresponding phrasing file. The split phrases are checked once again for note-to-note match before they are put into experiment.

4.4   Results

Six graduate students (not majored in music) were invited to record the samples. The number of mistakes they made are listed in Table 4.2.[3] These mistakes are identified using the unix `diff` [78] tool. Five of them (A to E) finished Clementi's entire Op.36, while performer F only recorded part of the work. The total number of recordings and the corresponding phrases/notes counts are shown in Table 4.3.

The number of phrases (according to our phrasing annotation) and notes are shown in Table 4.4. Fig. 4.1 shows the length distribution of each movement, most movements have around a few hundred notes, except the long No.6 and some short second movements. Fig. 4.2 shows the length distribution in numbers of phrases, most movements are around 20 phrases. The length distribution of the phrases in all six pieces are shown in Fig. 4.3, most phrases are shorter than 30 notes. Some very long phrases are usually virtuoso segments

---

[3]The performers are allowed to re-record as many time as they want, so the actual number of mistakes may be higher.

Table 4.2: Number of mistakes in the corpus, blank cell means the performer didn't record the piece

| Performer	1-1	1-2	1-3	2-1	2-2	2-3	3-1	3-2	3-3	4-1	4-2	4-3	5-1	5-2	5-3	6-1	6-2	Subtotal
A	0	5	2	4	3	0	4	2	2	4	5	9	9	2	3	4	1	59
B	2	1	1	2	2	1	6	0	3	2	3	6	12	3	3	10	7	64
C	1	1	0	1	0	1	2	0	0	3	2	3	10	1	35	6	1	67
D	0	1	1	2	3	1	4	1	1	10	6	3	10	2	7	13	2	67
E	2	3	4	4	0	3	4	0	0	21	6	22	23	3	9	18	13	135
F	1	3	2	11	6	8	7	2	6		15				20			81
Subtotal	6	14	10	24	14	14	27	5	12	40	37	43	64	11	77	51	24	473

Figure 4.1: Movement length (notes) distribution

with very fast note sequences, so it's hard to further split them.

Table 4.3: Total recorded phrases and notes count

Title	Recordings Count	Total Phrases	Total Notes
No.1 Mov. I	6	72	1332
No.1 Mov. II	6	60	882
No.1 Mov. III	6	102	1566
No.2 Mov. I	6	108	1920
No.2 Mov. II	6	36	750
No.2 Mov. III	6	168	2484
No.3 Mov. I	6	156	3156
No.3 Mov. II	6	42	444
No.3 Mov. III	6	120	2628
No.4 Mov. I	5	80	2325
No.4 Mov. II	6	78	1332
No.4 Mov. III	5	85	1920
No.5 Mov. I	5	85	3360
No.5 Mov. II	5	70	1580
No.5 Mov. III	6	144	3384
No.6 Mov. I	5	145	4180
No.6 Mov. II	6	78	2754
Total	97	1629	35997

Table 4.4: Phrases and notes count for Clementi's Sonatina Op.36

| Title	Phrases Count	Notes Count
No.1 Mov. I	12	222
No.1 Mov. II	10	147
No.1 Mov. III	16	261
No.2 Mov. I	18	320
No.2 Mov. II	6	125
No.2 Mov. III	28	414
No.3 Mov. I	25	526
No.3 Mov. II	6	74
No.3 Mov. III	19	438
No.4 Mov. I	25	465
No.4 Mov. II	12	222
No.4 Mov. III	16	384
No.5 Mov. I	17	672
No.5 Mov. II	13	316
No.5 Mov. III	24	564
No.6 Mov. I	28	836
No.6 Mov. II	11	459
Total	286	6445

Figure 4.2: Movement length (phrases) distribution

Figure 4.3: Phrase length (notes) distribution

Chapter 5

Experiments

In this chapter, we will show some experiment results to proof the effectiveness of our method. Section 5.1 deals with the onset deviation problem highlighted in Section 3.5.3. Section 5.2 discusses how various parameters in our system are chosen. Section 5.3 describes a subjective test to test if audience can or can't identify the difference between generated and human performances.

5.1   Onset Deviation Normalization

As mentioned in Section 3.5.3, a bad normalization method will usually result in unreasonable high onset deviation. To overcome this challenge, we proposed a automated way to select the normalization . In this section, we will evaluate the effectiveness of the method.

We extract the onset deviation feature from performer E's recording[1], using the two types of fixed normalization method and also the automatic normalization method mentioned in Section 3.5.3. The onset deviations extracted by each method are shown in Fig. 5.1, Fig. 5.2 and Fig. 5.3. Each dotted line from left to right represents a phrase in the corpus. Each dot represents the onset deviation value of a note. The notes are spread uniformly on the horizontal axis, which only shows the order of appearance, instead of the

---

[1]The effect of this method is less obvious for performer with better piano skill, because they have better control over tempo stability.

Figure 5.1: Onset deviations by aligning last note onset

real time scale. First, we can see in Fig. 5.2 that by aligning the note-off events of the last notes results in very large deviations in some phrases. This is because extending the last note in certain phrases to emphasize the ending is a common expression. This kind of extension will cause the last notes onset in the performance to be far apart from the score. Fig. 5.1 and Fig. 5.3 seemed to work better.Although they look similar, but the onset deviation values in Fig. 5.1 is more dramatic than those in Fig. 5.3, which proofs that the automatic normalization method can generally reduce the onset deviations. Another benefit of the automated normalization method over aligning last notes onset method is that the last notes are not force aligned, which allows more space for free expression for the last note. This effect can be seen in Fig. 5.1, in which the right-most end of a line, i.e. the last note, always goes back to zero, while in Fig. 5.3, the end of a line can end in different values

37

Figure 5.2: Onset deviations by aligning last notes note-off

Figure 5.3: Onset deviations using automated normalization method

5.2 Parameter Selection

5.2.1 SVM-HMM-related Parameters

There are many parameters which need adjustment in SVM-HMM. Two most important parameters, the termination accuracy ε and the misclassification penalty factor C in SVM, are systematically tested in this experiment to find the optimal value. Since SVM-HMM is an iterative algorithm, the ε parameter defines the required accuracy for the algorithm to terminate. A smaller ε will result in higher accuracy, but may take more iterations to compute. The C parameter determines how much weight should be assigned to penalise non-separable samples. A larger C will sacrifice larger margin for lower misclassification error, but it will make the execution time longer.

We split performer A's recordings into two sets: the training set includes pieces No.2 to No.6, and the testing set includes piece No.1. We train a model with the training set, and use the learned model to generate the testing set. The generated expressive performance is compared to the corresponding human recordings to calculate the accuracy of the prediction.

Ideally, the generated performance will be very similar in expression to the recording. In order to choose the best ε, we calculate the median of similarities between the generated and recorded performances for each ε choice. Note that each performance feature has its own model, so we will be looking at one performance feature and its ε parameter at a time. First, the generated performance feature sequence and the recorded one are normalized to a range from 0 to 1. This is because the generated performance may have the same up-and-downs as the score, but the value range may be different, so we use normalization to ease our these difference. The Euclidean distance between the two normalized sequences is calculated and divided by the length (number of notes) of the phrase, since the phrase can have arbitrary length. Similar procedure is applied to find the best C.

First we fixed C at 0.1 and tried different ε's: 100, 10, 1, 0.75, 0.5 and 0.1. Then, we fix ε at the optimal value determined in the previous step and test different C's: 10^{-3}, 10^{-2}, 10^{-1}, 0.5, 1, and 5. For each ε and C combination, we calculate the distance be-

Figure 5.4: Median distance between generated performances and recordings for different ε's

tween the generated pieces and recorded examples for all phrases in the testing set for each performer. Then we take the median of all these distances for each ε or C. The optimal ε or C is the one that minimize the median of the distances.

The median distance of the generated performance from the recording for various ε's are shown in Fig. 5.4. The execution time for various ε's are shown in Fig. 5.5. For ε value 100 and 10, the termination criteria is too generous so SVM-HMM terminates almost immediately without actually learned anything. Therefore, the outputs are a fixed value for any input. We abandon the data points for $\varepsilon = 100$ or 10. We can see that the distance drops slowly when ε becomes smaller. We choose $\varepsilon = 0.1$ for the best accuracy-time tradeoff.

As for different C parameter, the accuracy and execution time are shown in Fig. 5.6 and Fig. 5.7, respectively. We can't find a clear trend in Fig. 5.6, but we can find that for C over 10 and under 0.01, the model failed to produce meaning fule model (i.e. the output is a fixed value), so the data point is omitted in the figure. Therefore, choosing a

Figure 5.5: Execution time for different ε's

C in the middle will produce more robust model. In Fig. 5.7 the execution time grows as C goes larger, so considereing the robustness (always producing meaningful model) and time tradeoff, we choose C = 0.1 as our optimal C.

5.2.2 Quantization Parameter

Besides ε and C, the number of quantization levels for SVM-HMM input is also has some impact on the execution time. If the performance features are quantized into more fine-grained levels, the quantization errors can be reduced, but the execution time and memory usage will grow dramatically. Also, larger number of intervals doesn't imply more accurate or robust model. Because SVM-HMM is originally used in part-of-speech tagging problem, if we use divide the performance features into more intervals, there will be fewer samples in each interval. But from a statistical learning point of view, it is desirable to have fewer bins with more samples in each, rather than a large number of bins with very sparse samples in each. To illustrate this point, consider a three note segment is

Figure 5.6: Median distance between generated performances and recordings for different C's

Figure 5.7: Execution time for different C's

played once in the following MIDI velocity: (60, 70, 80), and the same segment is played again in (60.1, 69.9, 80.1). If we have a quantization interval width of, say, 0.05, then 60 and 60.1 may be quantized into different bins, and 70 and 69.9 may also be quantized into different bins, so the two phrases will be considered as two different case. However, if the quantization interval width is 1, both phrases may be quantized into the same label sequence, which is more desirable because the SVM-HMM algorithm can capture the similarity in the two samples.

Initially, we tried to quantized the values into 1025 uniform width bins, wishing to minimize the quantization error. But it take very long (hours, even days) to learn a model, and the output only falls on a very sparse set of values. So we reduce this number to 128. This level of quantization is fine enough to capture the performance nuance. Taking a rough estimate, onset deviation feature rarely exceeds ± 1, so the quantization interval width is around $\frac{1-(-1)}{128} = 0.015625$. Most duration ratios falls between zero and three, so the interval width is $\frac{3-0}{128} = 0.0234375$. MIDI velocity is roughly around 30 to 90, so the interval is about $\frac{90-30}{128} = 0.46875$. This level of granularity is good enough for our performance system, and can dramatically reduce the execution time without sacrificing the expressiveness of the models.

We repeated the ε selection experiment for quantization level of 1025 and 128. The execution time (in CPU second) is shown in Fig. 5.8. The time required for 1025 is larger than 128 by orders of magnitudes, but the expressiveness does not improve much.

5.3   Human-like Performance

The goal of our system is to create expressive, non-robotic music as oppose to deadpan MIDI. We will present two exprements to evaluate the performance of our system. First, we use the distance measure defined in Section 5.2 to see if the computer performance are more similar to human performance than inexpressive MIDI. Second, we need to perform a subjective test to verify if people can tell our generated performances apart from real human performances.

1518 expressive performance phrases were generated by our system. We follow a six-

Figure 5.8: Execution time for differnt number of quantization levels

fold cross-validation pattern: for each performer in the corpus, we use all his/her recorded phrases of Clementi's Op.36 No.2 to No.6 to train a model. Then the model is used to generate all phrases from Clementi's Op.36 No.1. The generate phrases and the performer's recordings of piece No.1 will all be included as samples. The process is repeated, but each time the piece excluded for training will be changed to No.2, No.3 and so on. So all six pieces will have a computer generated version (trained by each player's corpus) and a recorded version.

Ideally, by training the model by a certain performer's corpus, the generated performance will posses the same style as the performer. Therefore, the expression of the generated performance will be much similar to the performer's recordings, comparing to inexpressive MIDI score. Using the distance measure defined in Section 5.2, we can compare the 1518 generated performance with their corresponding human recording. Also the original score is transformed into MIDI and compared with the human recordings as control group. The average distance for different corpus-performance feature combination is

Table 5.1: Average (normalized) distance between generated performance and human recording, and between inexpressive MIDI and human performance

Performer		Duration	Onset	Loudness	Total
A	Generated-human	0.095	0.086	0.094	0.092
	MIDI-human	0.233	0.234	0.233	0.233
B	Generated-human	0.089	0.090	0.090	0.089
	MIDI-human	0.233	0.233	0.234	0.233
C	Generated-human	0.088	0.092	0.101	0.094
	MIDI-human	0.234	0.233	0.232	0.233
D	Generated-human	0.088	0.086	0.082	0.085
	MIDI-human	0.233	0.234	0.232	0.233
E	Generated-human	0.093	0.120	0.080	0.095
	MIDI-human	0.233	0.218	0.232	0.228
Total	Generated-human	0.091	0.097	0.088	0.092
	MIDI-human	0.233	0.228	0.233	0.231

listed in Table 5.1.

We can tell from Table 5.1 that a generated expressive performance are much similar to its corresponding human performance than original score MIDI. This result shows that our system is able to generate performance that are much similar to human the inexpressive MIDI.

Second, we performed a subjective survey to verify if the computer generated performance sounds better for human listener. In this survey, the same 1518 computer generated expressive phrases and their corresponding human recording were used as samples. Each test subject was given 10 randomly selected computer generated phrase and 10 human recordings, these 20 phrases are presented in random order. He/She was asked to rate each phrase according to the following criteria, which were proposed by the RenCon contest [1]:

1. Technical control: if a performance sounds like it is technically skilled thus performed with accurate and secure notes, rhythms, tempo and articulation.

2. Humanness: if the performance sounds like a human was playing it.

3. Musicality: how musical the performance is in terms of tone and color, phrasing, flow, mood and emotions

4. Expressive variation: how much expressive variation (versus deadpan) there is in the performance.

In RenCon, each judge was asked to give separate ratings for each criteria. But we believe this is too demanding for less-experienced participant, so we asked each test subject to give an overall rating from one to five. One being very bad, five being very good. The test subjects are also asked to report their musical proficiency in a three level scale:

1. No experience in music

2. Amateur performer

3. Professional musician, musicologist or student majored in music

We have also tried using all performers' recordings to train a single model. However, the expressive variation from that model is much smaller than a model trained by a single performer's recordings. This is because expression from different performers may cancel each other out. This phenomena can be found in the distribution histograms for each performance features (Fig. 5.9, Fig. 5.10 and Fig. 5.11). The features generated from the full corpus are slightly more concentrated, which results in less dramatic expression.

We received 119 valid samples for the survey. Fifty of them are from people with no music background, 59 are from amateur musicians, and the rest 10 are from professional musicians. The average rating given to computer generated performances and human recordings are listed in Table 5.2. It is clear that for professional and amateur musician, the average rating given to human performances are higher than computer performances. However, for participants who have no experience in music, the ratings are much closer. A Student T-test on the two ratings given by participants with no experience yields a p-value of 0.0312, therefore we can't reject the null hypothesis that the two ratings are different under a significance level of 99%. Therefore we can say for participants with no music experience, the computer generated music and human recordings are indistinguishable.

We also performed another survey for comparison: the test subjects are also given 20 samples, but the 10 computer generated samples are replaced by the the original score rendered as inexpressive MIDI. This survey received less feeback, only 27 valid suverys
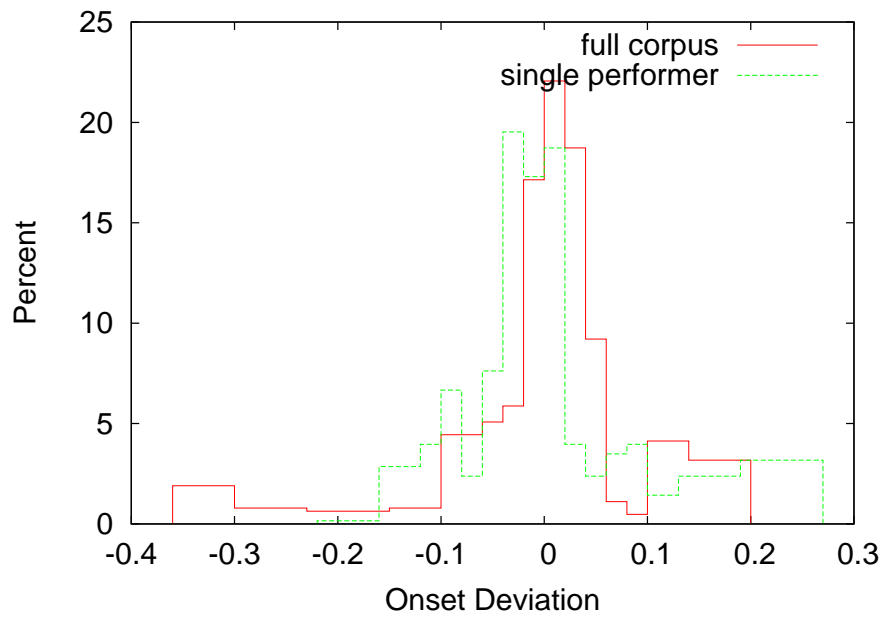
Figure 5.9: Distribution of onset deviation values from full corpus versus single performer's corpus
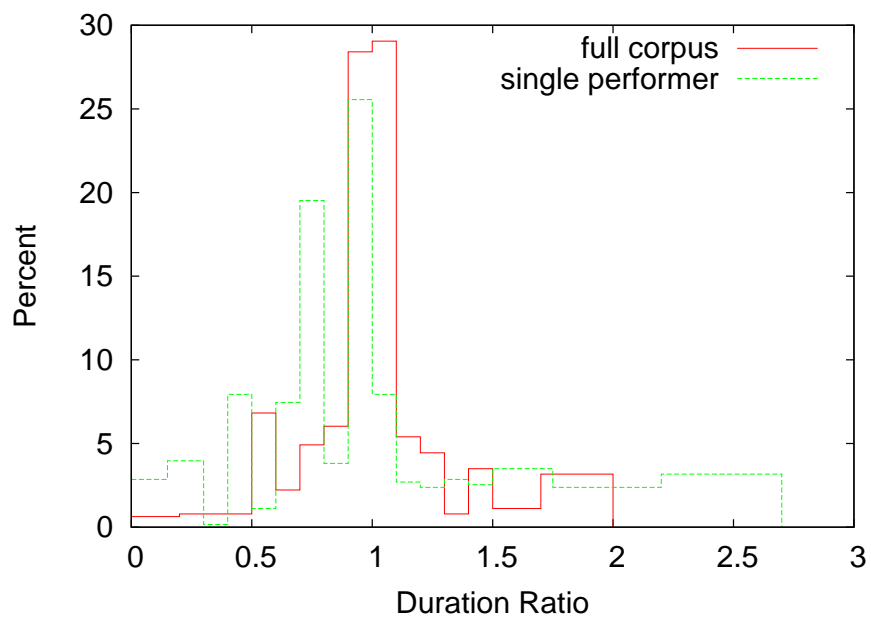
Figure 5.10: Distribution of duration ratio values from full corpus versus single performer's Corpus

Figure 5.11: Distribution of MIDI velocity values from full corpus versus single performer's corpus

Table 5.2: Average rating for generated performance and human recording; numbers in brackets are standard deviations

	Computer	Human
No experience	3.243 (1.036)	3.391 (0.986)
Amateur	2.798 (1.075)	3.289 (1.062)
Professional	2.430 (1.191)	3.010 (1.068)
Total	2.952 (1.102)	3.306 (1.035)

Table 5.3: Average ratings for inexpressive MIDI and human performance

	Human	MIDI
No experience	3.464 (1.011)	3.455 (1.089)
Amateur	3.046 (0.995)	3.369 (1.313)
Professional	2.633 (0.999)	2.333 (0.994)
Total	3.170 (1.035)	3.289 (1.237)

are collected, 11 are from people with no music background, 13 are from amateurs and 3 from professionals. The results are shown in Table 5.3.

From Table 5.2 and Table 5.3, we have the following findings: first, test subjects with no music background give very similar ratings to any kind of sample. Since they can't distinguish these samples from each other, we can simply ignore them. Amateurs and professional show much clear preference between samples. Amateurs prefer inexpressive MIDI over human performance, and they prefer human performance over computer performance. This results illustrates an important point: having human-like deviation is not the only factor that determins if a phrase sounds good. Since the timing and loudness deviation is very small, and the phrases are relatively short, it's not very easy to tell the difference between them, so any awkward expression will be the deterministic factor. Since the human performers we invited are not professionals, they may have performe clumsily in certain segments.Therefore, inexpressive MIDIs, which have no awkward expression, is rated higher than human performance. Since our system can't learn every nuance from training samples, the generated performnace may contain awkward or unnaturl expressions in addition to the already clumsy training samples, so the computer generated performances are rated even lower. Professionals can generally tell the difference between human performance and computer generated or inexpressive MIDI. But the ratings for computer generated performance and inexpressive MIDI do not have clear differences.

We can see from Table 5.2 and Table 5.3 that the standard deviation of the ratings are quite large, and our sample size is not large enough, so there are concerns to directly compare the average ratings. Instead, we can interpret the result in another apsect: if we look into each individual participant, we can check if a participant gives higher (average) rating to computer or human performances, or equal ratings for both. Similarly, we can

Table 5.4: Number of participants who gives higher rating to generated performance, human recordings or equal rating

	Computer	Equal	Human	Total
No experience	19	7	24	50
Amateur	7	3	49	59
Professional	1	1	8	10
Total	27	11	81	119

Table 5.5: Number of participants who gives higher rating to inexpressive MIDI, human recordings or equal rating

	MIDI	Equal	Human	Total
No experience	8	0	3	11
Amateur	9	0	4	13
Professional	1	0	2	3
Total	18	0	9	27

count the number of participants who rates inexpressive MIDI over human performanc or vice versa. The results are shown in Table 5.4 and Table 5.5, respectively. The results are similar to Table 5.2 and Table 5.3, but the difference are clearer.

Chapter 6

Conclusions

We have created a system that can perform monophonic score expressively. The expressive performance knowledge is learned from hum an recording using structural support vector machine with hidden Markov model output (SVM-HMM). We have also created a corpus consisting of scores and MIDI recordings. From subjective test, we find that the computer generated performance still can't achieve the same level of expressiveness of human performers. However, from our similarity measure, the computer generated expressive performance are much similar to the human performance than inexpressive MIDI.

There are many room for improvement. Structural expressions such as phrasing, contrast between sections, or even contrast between movements can be added, which requires automatic structural analysis. Other information like text notations, harmonic analysis and other musicological analysis can be added to the learning process. Supporting homophonic or polyphonic music is also important for the system to be useful. Sub-note expressions like physical model synthesizer or envelope shaping can also be applied to generate performances for specific musical instruments. It's also crucial to test the system on more samples of different genre or music style. We also believe that combining rule-based model and machine-learning model may be a possible direction for computer expressive music performance research. With rules serving as a high level guideline for structural expression, the machine-learning model can focus on note or sub-note level expression. User can gain more control by tweaking the rules.

Bibliography

[1] R. Hiraga, R. Bresin, K. Hirata, and H. Katayose, "Rencon 2004: Turing test for musical expression ," in *Proceedings of the 2004 conference on New Interfaces for Musical Expression (NIME '04)* (Y. Nagashima and M. Lyons, eds.), (Hamatsu, Japan), pp. 120--123, ACM Press, 2004.

[2] "Finale [Computer Software]." http://www.finalemusic.com/. [Online; accessed 2014-05-20].

[3] "Sibelius [Computer Software]." http://www.avid.com/us/products/sibelius/pc/Play-perform-and-share. [Online; accessed 2014-05-20].

[4] "Rachmianinoff - Plays Rachmaninoff [CD]." Zenph Music, 2009.

[5] "Cadenza [Computer Software]." http://www.sonation.net/. [Online; accessed 2014-05-20].

[6] A. Kirke and E. R. Miranda, "An Overview of Computer Systems for Expressive Music Performance," in *Guide to Computing for Expressive Music Performance* (A. Kirke and E. R. Miranda, eds.), pp. 1--47, Springer, 2013.

[7] A. Friberg, R. Bresin, and J. Sundberg, "Overview of the KTH rule system for musical performance," *Advances in Cognitive Psychology*, vol. 2, pp. 145--161, Jan. 2006.

[8] M. Hashida, N. Nagata, and H. Katayose, "Pop-E: a performance rendering system for the ensemble music that considered group expression," in *Proceedings of 9th In-*

ternational Conference on Music Perception and Cognition (M. Baroni, R. Addessi, R. Caterina, and M. Costa, eds.), (Bologna, Spain), pp. 526--534, ICMPC, 2006.

[9] S. R. Livingstone, R. Mühlberger, A. R. Brown, and A. Loch, "Controlling musical emotionality: an affective computational architecture for influencing musical emotions," *Digital Creativity*, vol. 18, pp. 43--53, Mar. 2007.

[10] N. P. M. Todd, "A computational model of rubato," *Contemporary Music Review*, vol. 3, pp. 69--88, Jan. 1989.

[11] N. P. McAngus Todd, "The dynamics of dynamics: A model of musical expression," *The Journal of the Acoustical Society of America*, vol. 91, p. 3540, June 1992.

[12] N. P. M. Todd, "The kinematics of musical expression," *The Journal of the Acoustical Society of America*, vol. 97, p. 1940, Mar. 1995.

[13] M. Clynes, "Generative principles of musical thought: Integration of microstructure with structure," *Journal For The Integrated Study of Artificial Intelligence*, 1986.

[14] M. Clynes, "Microstructural musical linguistics: composers' pulses are liked most by the best musicians," *Cognition*, 1995.

[15] M. Johnson, "Toward an expert system for expressive musical performance," *Computer*, vol. 24, pp. 30--34, July 1991.

[16] R. B. Dannenberg and I. Derenyi, "Combining instrument and performance models for high-quality music synthesis," *Journal of New Music Research*, vol. 27, pp. 211--238, Sept. 1998.

[17] R. B. Dannenberg, H. Pellerin, and I. Derenyi, "A Study of Trumpet Envelopes," in *Proceedings of the 1998 international computer music conference* (O. 1998, ed.), (Ann Arbor, Michigan), pp. 57--61, International Computer Music Association, 1998.

[18] G. Mazzola and O. Zahorka, "Tempo curves revisited: Hierarchies of performance fields," *Computer Music Journal*, vol. 18, no. 1, pp. 40--52, 1994.

[19] G. Mazzola, *The Topos of Music: Geometric Logic of Concepts, Theory, and Performance*. Basel/Boston: Birkhäuser, 2002.

[20] W. A. Sethares, *Tuning, Timbre, Spectrum, Scale*. Springer, 2005.

[21] H. Katayose, T. Fukuoka, K. Takami, and S. Inokuchi, "Expression extraction in virtuoso music performances," in *Proceedings of 10th International Conference on Pattern Recognition*, vol. I, pp. 780--784, IEEE Computer Society Press, 1990.

[22] H. Katayose, T. Fukuoka, K. Takami, and S. Inokuchi, "Extraction of expression parameters with multiple regression analysis," *Journal of Information Processing Society of Japan*, no. 38, pp. 1473--1481, 1997.

[23] O. Ishikawa, Y. Aono, H. Katayose, and S. Inokuchi, "Extraction of Musical Performance Rules Using a Modified Algorithm of Multiple Regression Analysis," in *International Computer Music Conference Proceedings*, (Berlin, Germany), pp. 348--351, International Computer Music Association, San Francisco, 2000.

[24] S. Canazza, G. De Poli, C. Drioli, A. Rodà, and A. Vidolin, "Audio Morphing Different Expressive Intentions for Multimedia Systems," *IEEE Multimedia*, vol. 7, pp. 79--83, July 2000.

[25] S. Canazza, A. Vidolin, G. De Poli, C. Drioli, and A. Rodà, "Expressive Morphing for Interactive Performance of Musical Scores," in *Proceedings of 1st International Conference on Web Delivering of Music*, p. 116, IEEE Computer Society, Nov. 2001.

[26] S. Canazza, G. De Poli, A. Rodà, and A. Vidolin, "An Abstract Control Space for Communication of Sensory Expressive Intentions in Music Performance," *Journal of New Music Research*, vol. 32, pp. 281--294, Sept. 2003.

[27] R. Bresin, "Artificial neural networks based models for automatic performance of musical scores," *Journal of New Music Research*, vol. 27, pp. 239--270, Sept. 1998.

[28] A. Camurri, R. Dillon, and A. Saron, "An experiment on analysis and synthesis of musical expressivity," in *Proceedings of 13th colloquium on musical informatics*, (L'Aquila, Italy), 2000.

[29] G. Grindlay, *Modeling expressive musical performance with Hidden Markov Models*. PhD thesis, University of Santa Cruz, CA, 2005.

[30] C. Raphael, "Can the computer learn to play music expressively?," in *Proceedings of the 8th International Workshop on Artificial Intelligence and Statistics* (T. Jaakkola and T. Richardson, eds.), pp. 113--120, Morgan Kaufmann, San Francisco, 2001.

[31] C. Raphael, "A Bayesian Network for Real-Time Musical Accompaniment.," *Neural Information Processing Systems*, no. 14, pp. 1433--1440, 2001.

[32] C. Raphael, "Orchestra in a box: A system for real-time musical accompaniment," in *Proceedings of 2003 International Joint conference on Artifical Intelligence (Working Notes of IJCAI-03 Rencon Workshop)* (G. Gottob and T. Walsh, eds.), (Acapulco, Mexico), pp. 5--10, Morgan Kaufmann, San Francisco, 2003.

[33] L. Dorard, D. Hardoon, and J. Shawe-Taylor, "Can style be learned? A machine learning approach towards 'performing' as famous pianists.," in *Proceedings of the Music, Brain and Cognition Workshop -- Neural Information Processing Systems*, Whistler, Canada, 2007.

[34] M. Wright and E. Berdahl, "Towards machine learning of expressive microtiming in Brazilian drumming," in *Proceedings of the 2006 International Computer Music Conference* (I. Zannos, ed.), (New Orleans, USA), pp. 572--575, ICMA, San Francisco, 2006.

[35] R. Ramirez and A. Hazan, "Modeling Expressive Music Performance in Jazz.," in *Proceedings of 18th international Florida Artificial Intelligence Research Society Sonference (AI in Music and Art)*, (Clearwater Beach, FL, USA), pp. 86--91, AAAI Press, Menlo Park, 2005.

[36] R. Ramirez and A. Hazan, "Inducing a generative expressive performance model using a sequential-covering genetic algorithm," in *Proceedings of 2007 annual conference on Genetic and evolutionary computation*, (London, UK), ACM Press, New York, 2007.

[37] Q. Zhang and E. Miranda, "Towards an evolution model of expressive music performance," in *Proceedings of the 6th International Conference on Intelligent Systems Design and Applications* (Y. Chen and A. Abraham, eds.), (Jinan, China), pp. 1189--1194, IEEE Computer Society, Washington, DC, 2006.

[38] E. Miranda, A. Kirke, and Q. Zhang, "Artificial evolution of expressive performance of music: An imitative multi-agent systems approach," *Computer Music Journal*, vol. 34, no. 1, pp. 80--96, 2010.

[39] Q. Zhang and E. R. Miranda, "Evolving Expressive Music Performance through Interaction of Artificial Agent Performers," in *Proceedings of ECAL 2007 workshop on music and artificial life (MusicAL 2007)*, (Lisbon, Portugal), 2007.

[40] J. L. Arcos, R. L. De Mántaras, and X. Serra, "SaxEx: a case-based reasoning system for generating expressive musical performances," in *Proceedings of 1997 International Computer Music Conference* (P. Cook, ed.), (Thessalonikia, Greece), pp. 329--336, ICMA, San Francisco, 1997.

[41] J. L. Arcos, R. L. De Mántaras, and X. Serra, "SaxEx: A case-based reasoning system for generating expressive musical performances," *Journal of New Music Research*, vol. 27, no. 3, pp. 194--210, 1998.

[42] J. L. Arcos and R. L. De Mántaras, "An Interactive Case-Based Reasoning Approach for Generating Expressive Music," *Journal of Applied Intelligence*, vol. 14, pp. 115--129, Jan. 2001.

[43] T. Suzuki, T. Tokunaga, and H. Tanaka, "A case based approach to the generation of musical expression," in *Proceedings of the 16th International Joint Conference on*

Artificial Intelligence, (Stockholm, Sweden), pp. 642--648, Morgan Kaufmann, San Francisco, 1999.

[44] T. Suzuki, "Kagurame phase-II," in *Proceedings of 2003 International Joint Conference on Artificial Intelligence (working Notes of RenCon Workshop)* (G. Gottlob and T. Walsh, eds.), (Acapulco, Mexico), Morgan Kaufmann, Los Altos, 2003.

[45] K. Hirata and R. Hiraga, "Ha-Hi-Hun: Performance rendering system of high controllability," in *Proceedings of the ICAD 2002 Rencon Workshop on performance rendering systems*, (Kyoto, Japan), pp. 40--46, 2002.

[46] G. Widmer, "Large-scale Induction of Expressive Performance Rules: First Quantitative Results," in *Proceedings of the 2000 International Computer Music Conference* (I. Zannos, ed.), (Berlin, Germany), pp. 344--347, International Computer Music Association, San Francisco, 2000.

[47] G. Widmer and A. Tobudic, "Machine discoveries: A few simple, robust local expression principles," *Journal of New Music Research*, vol. 32, pp. 259--268, 2002.

[48] G. Widmer, "Discovering simple rules in complex data: A meta-learning algorithm and some surprising musical discoveries," *Artificial Intelligence*, vol. 146, pp. 129--148, 2003.

[49] G. Widmer and A. Tobudic, "Playing Mozart by Analogy: Learning Multi-level Timing and Dynamics Strategies," *Journal of New Music Research*, vol. 32, pp. 259--268, Sept. 2003.

[50] A. Tobudic and G. Widmer, "Relational IBL in music with a new structural similarity measure," in *Proceedings of the 13th International Conference on Inductive Logic Programming* (T. Horvath and A. Yamamoto, eds.), pp. 365--382, Springer Verlag, Berlin, 2003.

[51] A. Tobudic and G. Widmer, "Learning to play Mozart: Recent improvements," in *Proceedings of 2003 International Joint conference on Artifical Intelligence (Working Notes of IJCAI-03 Rencon Workshop)* (K. Hirata, ed.), (Acapulco, Mexico), 2003.

[52] P. Dahlstedt, "Autonomous evolution of complete piano pieces and performances," in *Proceedings of ECAL 2007 workshop on music and artificial life (Music AL 2007)*, (Lisbon, Portugal), 2007.

[53] A. Kirke and E. Miranda, "Using a biophysically-constrained multi-agent system to combine expressive performance with algorithmic composition," 2008.

[54] L. Carlson, A. Nordmark, and R. Wikilander, *Reason version 2.5 -- Getting Started*. Propellerhead Software, 2003.

[55] Y.-H. Kuo, W.-C. Chang, T.-M. Wang, and A. W. Su, "TELPC BASED RE-SYNTHESIS METHOD FOR ISOLATED NOTES OF POLYPHONIC INSTRU-MENTAL MUSIC RECORDINGS," in *Proceedings of the 16th International Conference on Digital Audio Effects (DAFx-13)*, (Maynooth, Ireland), pp. 1--6, 2013.

[56] T. Joachims, T. Finley, and C.-N. J. Yu, "Cutting-plane training of structural SVMs," *Machine Learning*, vol. 77, pp. 27--59, May 2009.

[57] I. Tsochantaridis, T. Joachims, T. Hofmann, and Y. Altun, "Large Margin Methods for Structured and Interdependent Output Variables," *Journal of Machine Learning Research*, vol. 6, pp. 1453--1484, 2005.

[58] Y. Altun, I. Tsochantaridis, and T. Hofmann, "Hidden Markov Support Vector Machines," in *Proceedings of the 20th International Conference on Machine Learning*, vol. 3, (Washington DC, USA), pp. 3--10, 2003.

[59] M. Cuthbert and C. Ariza, "music21 [computer software]." http://web.mit.edu/music21/. [accessed 2014-05-20].

[60] MIDI Manufacturers Association, "The Complete MIDI 1.0 Detailed Specification." http://www.midi.org/techspecs/midispec.php. [Online; accessed 2014-05-20].

[61] S. H. Lyu and S.-K. Jeng, "COMPUTER EXPRESSIVE MUSIC PERFORMANCE BY PHRASE-WISE MODELING," in *Workshop on Computer Music and Audio Technology*, 2012.

[62] T. Joachims, "SVMˆhmm: Sequence Tagging with Structural Support Vector Machines." http://www.cs.cornell.edu/people/tj/svm_light/svm_hmm.html. [Online; accessed 2014-05-20].

[63] "MusicXML 3.0 Specification." http://www.musicxml.com/for-developers/. [Online; accessed 2014-05-20].

[64] R. P. Brent, *Algorithms for Minimization Without Derivatives*. 2013.

[65] M. Hashida, T. Matsui, and H. Katayose, "A New Music Database Describing Deviation Information of Performance Expressions," in *International Conference of Music Information Retrival (ISMIR)*, pp. 489--494, 2008.

[66] S. Flossmann, W. Goebl, M. Grachten, B. Niedermayer, and G. Widmer, "The Magaloff project: An interim report," *Journal of New Music Research*, vol. 39, no. 4, pp. 363--377, 2010.

[67] W. Goebl, S. Flossmann, and G. Widmer, "Computational investigations into between-hand synchronization in piano playing: Magaloff's complete Chopin," in *Proceedings of the Sixth Sound and Music Computing Conference*, pp. 291--296, 2009.

[68] M. Grachten and G. Widmer, "Explaining musical expression as a mixture of basis functions," in *Proceedings of the 8th Sound and Music Computing Conference (SMC 2011)*, 2011.

[69] S. Flossmann, W. Goebl, and G. Widmer, "Maintaining skill across the life span: Magaloff's entire Chopin at age 77," in *Proceedings of the International Symposium on Performance Science*, 2009.

[70] M. Grachten and G. Widmer, "Linear basis models for prediction and analysis of musical expression," *Journal of New Music Research*, 2012.

[71] S. Flossmann, M. Grachten, and G. Widmer, "Expressive performance rendering with probabilistic models," in *Guide to Computing for Expressive Music Performance* (A. Kirke and E. R. Miranda, eds.), pp. 75--98, Springer London, 2013.

[72] S. Flossman and G. Widmer, "Toward a model of performance errors: A qualitative review of Magaloff's Chopin," in *International Symposium on Performance Science*, (Utrecht), AEC, 2011.

[73] S. Flossmann, W. Goebl, and G. Widmer, "The Magaloff corpus: An empirical error study," in *Proceedings of the 11th ICMPC*, (Seattle, Washington, USA), 2010.

[74] G. Widmer, S. Flossmann, and M. Grachten, "YQX Plays Chopin," *AI Magazine*, vol. 30, p. 35, July 2009.

[75] F. Lerdahl and R. S. Jackendoff, *A Generative Theory of Tonal Music*. 1983.

[76] "KernScores." http://kern.ccarh.org/. [Online; accessed 2014-05-20].

[77] M. Clementi, *SONATINES pour Piano a 2 mains Op. 36 VOLUME I [Musical Score]*. Paris: Durand & Cie., plate d. & c. 9318 ed., 1915.

[78] P. Eggert, M. Haertel, D. Hayes, R. Stallman, and L. Tower, "diff [Computer Program]."

[79] M. Good, "MusicXML: An Internet-Friendly Format for Sheet Music," in *XML Conference hosted by IDEAlliance*, 2001.

[80] E. Selfridge-Field, *Beyond MIDI: The Handbook of Musical Codes*. MIT Press, 1997.

[81] "LilyPond." http://www.lilypond.org. [Online; accessed 2014-05-20].

Appendix A

Software Tools Used in This Research

This research won't come into reality without many free and open-source software tools and free resources, we will walk you through a brief introduction to the softwares we used in this research.

Linux Operating System

Most of the tools introduced below runs on modern Linux distributions. The distribution we are using is **Linux Mint Debian Edition (LMDE)**[1] (Linux kernel 3.10), which is a user-friendly Linux distribution based on Debian Testing. User who want to try music-related softwares without installing Linux on their harddrive can try **64 Studio**[2] Linux, which is a live CD distribution with many music-related software pre-installed. It also has many kernel optimizations for real-time music manipulation. **Ubuntu Studio**[3] is also an option, which has many pre-installed music softwares and is based on the popular Ubuntu Linux.

Many Linux distributions use PulseAudio audio server to manage audio device. But a badly configured PulseAudio server will introduce severe latency, which is not acceptable while doing MIDI recording. One workaround is to remove PulseAudio and use raw ALSA (Advanced Linux Sound Architecture) driver instead. But be careful, hardware

---

[1] http://www.linuxmint.com/download_lmde.php
[2] http://www.64studio.com/
[3] http://ubuntustudio.org/

volume keys may not work without PulseAudio.

Programming Languages

Python

Many researcher will choose Matlab or Octave for scientific projects because they have many useful toolboxes included. However, we believe that research project "doesn't exist in vacuum". Drawing insight from the famous 80-20 rule, only 20% of the code are actually doing the core algorithm, the rest 80% are doing file manipulation, configuration, user interaction, and visualization. Therefore, choosing a powerful and easy to write general-purpose programming language is extremely crucial. **Python**[4] construct most of the infrastructure code for this project. Python is super easy to code, and has almost every tool you need to construct a fully functional experiment environment. We will highlight some useful module:

`Music21`[5]

We would like to give special thanks to the `music21` developemnt team. `Music21` is a Python toolbox for music notation manipulation and analysis, developed by MIT. `Music21` can parse many score notations like MusicXML, MIDI[6] and more into a very convenient `music21` object data structure. Researcher can easily filter, split, search, and transform music notations. There are also many music analysis routines and feature extractors included. If you want to do computer music research, `music21` is a god-sent resource.

---

[4]https://www.python.org/
[5]http://web.mit.edu/music21/
[6]By default, `music21` will quantize MIDI input, so if you want to import MIDI recorded from human performance, you need to bypass the default parser and manually disable the quantizer.

SciPy, NumPy and Matplotlib[7]

SciPy is a project that contains many useful toolboxes for scientific computation in Python. The SciPy core library and NumPy provides numerical and vector calculation for Python, with similar capability to Matlab. Matplotlib provides plotting tools also similar to Matlab. It's useful for small scale calculation, but heavy duty mathematical calculation, we suggest `R` programming language, which will be discussed in later section.

`Simplejson`

JSON (JavaScript Object Notation) is a plaintext data-interchange format, similar to XML but much light-weight. JSON is useful in experiment code for two purpose: first, JSON can serve as configuration file, it easy to parse and easy to edit. Second, JSON can serve as intermediate data file between each experiment module. For example, we use JSON to send extracted features from feature extractors to the machine learning module. Although plaintext takes more storage than binary file, but it's much easier to debug because it's human readable. And you can simply parse the intermediate values and plot it using other plotting program. Python provides build-in support for JSON format via `json` and `simplejson` packages.

`Argparse`

`Argparse` provides command line argument parser for Python scripts, using commandline arguments with configuration file in JSON, you can create very flexible, extendible scripts that are easy to automate.

`Logging`

The built in `logging` module can print logging information with predefined format, and it supports log level. By using log level, you can print debug information during development, and hide all debug message during production simply by changing the log

---

[7]http://www.scipy.org/

64

level flag.

R[8]

R is a programming language for statistical calculation, but it can also do general purpose math and plotting very well. R follows a functional programming design, so it may take some time to learn for people who only have experience in C/C++, Java or other imperative and/or Object-oriented programming language. But it is a great tool for statistical computation, data analysis and visualization. We use R for experiment data analysis and for linear regression in early version of this research. R and Python can work seamlessly through the `rpy` package.

Score Manipulation and Corpora

MusicXML and MuseScore

MusicXML[9] is a digital score notation format based on XML. It is well supported in most commercial music typesetting software. To view and edit musicXML score, we use the open-source software **MuseScore**[10], it provides basic editing capability, and can export score as PDF. However, MuseScore often crash while loading bad-formatted musicXML file, so sometimes you need to look into it log file and fix the ill-formated XML via a text editor.

Corpora

`Music21` contains a corpus[11], which will be automatically installed if you accept the licence term during `music21` installation. It covers a wide range of composers from early music, classical music to folk songs, with various genre and musical style. Another public

---

[8] http://www.r-project.org/
[9] http://www.musicxml.com/
[10] http://musescore.org/
[11] http://web.mit.edu/music21/doc/systemReference/referenceCorpus.html

available corpus is called **KernScore**[12], which provides a better search engine. You can find works by composer, genre, form or other criteria. There are even a special section containing monophonic works. Scores from both corpus can be loaded and transformed in to desired format via `music21`.

MIDI Recording

Rosegarden[13] is a digital audio workstation (DAW) software designed specifically for MIDI. It can record, edit, mix and export MIDI tracks. To actually hear the music, you need a MIDI synthesizer to work with Rosegarden. **Timidity++**[14] is built-in in many Linux distribution, and it provides a commandline interface to synthesize MIDI directly into a WAV file. However, the default sound quality from Timidity++ is not very satisfying, so we suggest qSynth, which is a QT front end for **FluidSynth**[15]. The default soundfont that comes with FludiSynth has very good sound quality.

With all these music software, it will soon be very hard to control the interconnection between programs. This is when **JACK**[16] comes to help. JACK is like a virtual "plug-board" for software that implements the JACK interface. It provides a central place in which you can control how the music data flows between programs and hardware.

Audio Manipulation

When MIDI files are synthesized into WAV format, there are many tools that can help editing them. The most easy to use software with GUI is **Audacity**[17], it can edit and mix audio tracks. For commandline tools (in case you need automation), **oggenc**[18](ogg

---

[12]http://kern.ccarh.org/
[13]http://www.rosegardenmusic.com/
[14]http://timidity.sourceforge.net/
[15]http://sourceforge.net/projects/fluidsynth/
[16]http://jackaudio.org/
[17]http://audacity.sourceforge.net/
[18]http://www.vorbis.com/

vorbis encoder), **lame**[19](MP3[20] encoder) and **FFmpeg**[21] are very helpful for file format transformation. To cut and combine audio tracks from commandline, use **SoX**[22].

Data Visualization

As mentioned before, R and Matplotlib are good candidate for visualizing experiment data. But if you don't want to learn the syntax of R or Python, you can try **gnuplot**[23]. Gnuplot is a interactive (and scriptting) environment for generating various types of plot like line plots or bar charts. It works particularly well if you use grep to extract data for many files, say, extracting execution time information from logs.

SVM^{hmm}

SVM[hmm24] is an implementation for structural support vector machine with hidden Markov model output. It's developed by Thorsten Joachims from Cornell University. It is based on SVM^{struct}, a more general framework for structural support vector machine. There are many other SVM^{struct} extensions such as Python or Matlab API.

Other

Sometimes the machine learning algorithm will run for a very long time. Then it's better if you can find a server that runs 24-7 in your home or laboratory. You can install a **ssh** server on that machine, and controls the experiment execution remotely. However, the experiment program will be terminated once you log out the ssh session. You can run your experiment program in **tmux**[25], a terminal multiplexer, instead. It will keep your program running even if you log out of your SSH session.

---

[19]http://lame.sourceforge.net/
[20]Please consider open format like ogg first, MP3 is a closed format and may have patent issues.
[21]http://www.ffmpeg.org/
[22]http://sox.sourceforge.net/
[23]http://www.gnuplot.info/
[24]http://www.cs.cornell.edu/people/tj/svm_light/svm_hmm.html
[25]http://tmux.sourceforge.net/

Modern machines often have multi-core CPUs. But if your program only runs in one core, you waste the CPU resources and also your time. **Gnu-parallel**[26] can dispatch multiple instances of your script or program to each core. It will automatically find new job to run when the previous one is finished, so the CPU will always run on its full capacity.

Finally, We use **git**[27] for version control (including code and document). And LaTeX[28] is used to typeset this document.

Summary

We have reviewed many software tools used to construct this research. We want to emphasize that it is totally possible to use *only* free and open-source software to do all these heavy lifting. We encourge the reader to try these tools out, spread the words and even contribute to these projects. By doing so we can create a more friendly scientific computing community and make the world a better place.

---

[26]http://www.gnu.org/software/parallel/
[27]http://git-scm.com/
[28]http://latex-project.org/